

Intro to Parallel Programming

Nate Hattersley, Justin Latona

Fall 2024




The Modern Computer


- ▶ Modern computers have multiple cores
- ▶ Can process multiple instructions simultaneously
- ▶ Lets you watch football while doing homework
- ▶ We can use multiple cores to speed up estimation procedures
- ▶ Broadly known as parallel computing

Models of Parallel Computation I

Two important concepts are *threads* and *processes*



Multi-Threading



Multi-processing

Multiple threads can run inside one parent process. Operating system schedules processes and threads across available cores

Models of Parallel Computation II

- ▶ Why prefer multi-processing over multi-threading? Data races. A Julia example:

```
using Base.Threads

winner = 0
@threads for i in 1:10
    winner = i
end
@show winner
```

- ▶ `winner` is indeterminate. What would happen in mutli-processing?

Models of Parallel Computation III

- ▶ A third model to understand is *cooperative multitasking*
- ▶ When making API calls, your computer is just waiting while the server responds
- ▶ Why not fire off more requests while waiting for the response?
- ▶ Single *thread*, multiple *tasks*
- ▶ For IO-bound applications (whereas multi-threading/processing is for CPU-bound)
- ▶ See `asyncio` and `aiohttp` in Python for the curious

Loops vs. Implicit Parallelization I

- Implicit parallelization: when programming language allows compiler or interpreter to automatically exploit parallelism inherent to a problem through the language's constructs.

```
import numpy as np

x = np.random.normal(0,1,10**7) # 10 million random draws
y = np.square(x)                # takes 0.05 seconds on my PC
y = x**2                        # about the same;
                                # numpy overrides standard math operators,
                                # invokes implicitly parallelized functions

y = np.zeros(x.shape[0])
for i in range(x.shape[0]):
    y[i] = x[i]**2               # takes 2.96 seconds on my PC
```

- Avoid loops wherever possible. Leverage implicit parallelization with vectorization and broadcasting.

Loops vs. Implicit Parallelization II

- ▶ “Vectorization” in high-level language (e.g. python, matlab, julia) means use of optimized, pre-compiled code in low-level language (e.g. C) to perform mathematical operations over sequence of data.
- ▶ E.g. numpy takes advantage of homogenous dtype in an array to delegate task of performing mathematical operations on array’s contents to pre-compiled C code.
- ▶ Examples: square (unary); np.maximum (binary); np.sum (sequential)

```
import numpy as np

x = np.random.normal(0,1,10**7) # 10 million random draws
y = np.sum(x)                   # takes 0.02 seconds on my PC

y = 0
for i in range(x.shape[0]):
    y += x[i]                   # takes 1.91 seconds on my PC
```

- ▶ Most numpy math functions are “vectorized,” including logical and LA functions.

Broadcasting I

- ▶ Mechanism allowing arrays of unequal shapes to be used in vectorized operations.
- ▶ Effectively “stretches” arrays to replicate contents along chosen dimensions, such that higher-dim array suits the operation.
- ▶ Enables element-wise calculation across otherwise incompatible shapes.
- ▶ E.g. let's raise a length-10 vector to powers (element-wise) contained in the columns of a (10, 10 million) matrix.

Broadcasting II

```
import numpy as np

X = np.random.uniform(1,2,size=(10,10**7))
Y = np.random.uniform(2,10,size=10)

Z = Y**X                                # ValueError: operands could not be broadcast
                                         # together with shapes (10,) (10,10000000)

Z = np.zeros((X.shape))
for i in range(X.shape[1]):              # could do it this way, takes 43.5 seconds
    Z[:,i] = Y**X[:,i]

Z = Y[:,None]**X                          # broadcasting is MUCH faster: 8.9 seconds
```

- ▶ Gains are magnified over larger loops and more complicated calculations.
- ▶ Might have many such instances in a single structural model. Utilizing implicit parallelization in place of loops can make tractable an otherwise intractable model.

Julia

```
using Distributed; addprocs(4)

@everywhere begin
    using Statistics
    a = 4
end

pmap(1:100) do i
    mean(randn(a)) * i
end
```

R

```
library(parallel)

cl <- makePSOCKCluster(4)
a <- 4

clusterExport(cl, "a")

parSapply(cl, 1:100, function(i) mean(rnorm(a)) * i)
```

Python

```
import numpy as np, functools
from multiprocessing import Pool

def par_func(rng,a,i):
    return np.mean(rng.normal(size=a))*i

def outer(n_procs):
    rng = np.random.default_rng(1)
    a = 4
    with Pool(n_procs) as p:
        p.map(functools.partial(par_func,rng,a),np.arange(1,100).tolist())

if __name__ == '__main__':
    outer(4)
```

- ▶ Avoid referring to globally defined objects, pass all function inputs explicitly to child processes (as in Julia and R) by freezing into function signature.
- ▶ This is “best practice” for pseudo-random number generation in `numpy`, but *not* when the PRNGs are being drawn in parallel. More on this later.

Matlab

```
clear; clc;  
  
a = 4;  
P = 4;  
pool = parpool('local', P);  
  
parfor i = 1:100  
    mean(randn(a,1)) * i  
end  
  
delete(pool);
```

Example: Multinomial Choice Problem - Setup

- ▶ 100,000 consumers i , each with an observed scalar individual characteristic W_i .
- ▶ 10 product choices j , each with scalar product characteristic X_j .
- ▶ Unobserved consumer-specific component of indirect utility S_i is i.i.d. std. normal
- ▶ Use 1,000 simulation draws S_{is} for each consumer to form a simulated likelihood.
- ▶ Assuming parameters $\beta = (\beta_0, \beta_1, \beta_2, \beta_3)$, indirect utility for i -th consumer, j -th product and s -th simulant is

$$U_{ijs} = \beta_0 + (\beta_1 + \beta_2 W_i + \beta_3 S_{is}) X_j + \epsilon_{ijs}$$

- ▶ ϵ_{ijs} is idiosyncratic unobserved utility, assumed distributed type-1 extreme value.
- ▶ Draw (X_j, W_i, S_{is}) , generate fake consumer choices and parameters.
- ▶ Parallelize formation of likelihood over consumer-choice pairs, integrating-out simulants to compute CCPs.

Multinomial Choice (python + multiprocessing): DGP

```
import numpy as np
from multiprocessing import Pool
import functools

def data():
    N_cons = 100000          # 100,000 consumers (W)
    N_choices = 10           # 10 choices (X)
    N_sims = 1000            # 1000 sim draws (S)
    W = np.random.uniform(0,1,N_cons)
    X = np.random.uniform(0,1,N_choices)[:,None]
    S = np.random.normal(0,1,N_sims)[None,:] # prep for broadcasting
    Y = np.random.randint(0,10,N_cons) + 1   # 100k consumer choices (Y)
    prod_IDs = np.arange(10).astype(int) + 1 # product IDs 1-10
                                           # match choices to prod. IDs
                                           # to create binary choice matrix
                                           # (N_cons x N_prods)

    aux = Y[:,None] / prod_IDs[None,:]
    choices = np.zeros((aux.shape[0],aux.shape[1]))
    choices[aux==1] = 1
    beta = np.random.uniform(0,1,4)          # fake parameters
    W_list = W.tolist()                     # send W to list for p.map
    return W_list, X, S, choices, beta
```

Multinomial Choice (python + multiprocessing): Compute Likelihood

```
def comp_CCP(X, S, beta, W_list):
    e_util_ijs = np.exp(beta[0] + (beta[1] + beta[2]*W_list + beta[3]*S)*X)
    CCP_ijs = e_util_ijs/np.sum(e_util_ijs, axis=0) # compute CCPs after broadcast
    CCP_ij = np.mean(CCP_ijs, axis=1) # integrate-out simulants
    return CCP_ij # length-10 vector for i

def total_LL(n_processes):
    W_list, X, S, choices, beta = data()
    # freeze inputs into CCP function signature
    partial_comp_CCP = functools.partial(comp_CCP, X, S, beta)
    with Pool(n_processes) as p: # call process pool
        # p.map takes iterable of consumer characteristics
        # and generates list of CCP function outputs
        CCP_ij = np.vstack(p.map(partial_comp_CCP, W_list))
        total_LL = np.sum(choices*np.log(CCP_ij) + (1-choices)*np.log(1-CCP_ij))
    return total_LL

# need to declare main module, any calls above
if __name__ == '__main__': # are sent to child processes and duplicated
    n_processes = 20 # allocate 20 cores (of 40 available on VM)
    tot_LL = total_LL(n_processes)
```


Multinomial Choice (python): Discussion I

- ▶ Runs in 1.1 seconds on 20 cores.
- ▶ “Serial” implementation (also iterating over consumers, but “mapped”) takes about 11 seconds.
- ▶ Or, vectorize/broadcast to avoid the loop over consumers? Work in 3D:

```
W = np.random.uniform(0,1,N_cons)[:,None,None]
X = np.random.uniform(0,1,N_choices)[None,:,None]
S = np.random.normal(0,1,N_sims)[None,None,:]
# other corresponding modifications to functions...
# (see Github repository for code)
```

- ▶ Takes about 30 seconds. Why? We said broadcasting was faster than a loop.
- ▶ The issue is `np.exp` of 3D tensor. Only 10s of operation is broadcast to produce 3D tensor, other 20s is exponential of it.

Multinomial Choice (python): Discussion II

- ▶ Some numpy functions (powers especially) don't vectorize well into high-dimensions. Lots of overhead.

- ▶ Numerical example:

```
np.exp time over 100 runs, 10^6 draws: 10.19 ms  
mapped math.exp time over 100 runs, 10^6 draws: 0.00026 ms  
for-loop math.exp time over 100 runs, 10^6 draws: 339.47 ms
```

- ▶ `map` is 39,000 times faster than `np.exp`; 1.3 million times faster than true serial implementation.
 - ▶ `map` "works as an iterator to return a result after applying a function to every item of an iterable."
 - ▶ Still iterating, but effectively manually vectorizes function onto iterable.
- ▶ `np.exp` is implicitly parallel but still slower.
- ▶ A true serial (for-loop) implementation over consumer features in MNC problem would take too long, haven't computed it.

Multinomial Choice (python): Discussion III

- ▶ Broadcasting works well in high-dimensions with matrix algebra:
 - ▶ `matmul`,
 - ▶ broadcast operations over arrays, etc.
 - ▶ Depending on operation, should try different approaches.
- ▶ Taking exponents of a 3D matrix is inherent challenge to computing CCPs (per guess of params) in estimating RC logit demand models with product features and both observed and unobserved individual heterogeneity.

Multinomial Choice (python): Discussion IV

- ▶ Workarounds include:
- ▶ Iterating over a dimension and taking exponentials of 2D matrices. This should be further optimized:
 - ▶ Use `map` to “manually vectorize” (as above)
 - ▶ `flatten` a dimension + `bincount` → 2D, no loop
- ▶ And/or “just-in-time” compilation: native machine level code optimized to your functions (compile to GPU or CPU); what Julia + Matlab do automatically. Powerful.
- ▶ In Python: `jax` and `numba` packages.
- ▶ `jax.jit` compiling the CCP and likelihood calc. over 3D takes 3.8 seconds (vs. 30).
- ▶ Good idea to use `jit`-compilation in large problems. Need to try `jit` with `map` to further optimize MNC example.

Multinomial Choice (python): Discussion V

- ▶ In many situations, you would not be able to replace your entire parallelized operation with broadcasting anyway; need a loop.
- ▶ Gains from parallelization can be even larger in real-world problems, depending on what you're parallelizing.
- ▶ E.g. bootstrap; iterating over estimation starting points; simulations; running regressions on many separate datasets.
- ▶ Key is the operation needs to be "parallelizable:" Passing *many* objects separately through one process (process can change given the iterable).
- ▶ Process pool can handle multiple iterables: `Pool.starmap` in python, or package into a list of tuples and use `Pool.map`. Read more, [here \(link\)](#).

Resource Monitoring

- ▶ Windows: Ctrl+Alt+Del (Task Manager); Mac: Activity Monitor
- ▶ Mac, Linux and Cygwin Terminals: `top`/`htop`

```

0[|||||] 59.1% 4[|||||] 27.5%
1[|||||] 84.4% 5[|||||] 22.7%
2[|||||] 90.7% 6[|||||] 14.2%
3[|||||] 49.3% 7[|||||] 10.7%

Mem[|||||] 3.006/8.006G Tasks: 398, 1603 thr, 0 kthrs; 2 running
Swp[|||||] 1.116/2.006G Load average: 6.00 5.13 4.19
Uptime: 02:44:58

Main
PID USER PRI NI VIRT RES S CPU% MEM% TIME+ Command
1102 nate 17 0 4236 55888 ? 44.5 0.7 55:00.00 /Applications/Firefox Nightly.app/Contents/MacOS/media-plugin-r
985 nate 24 0 4050 355M ? 33.0 4.3 54:12.00 /Applications/Firefox Nightly.app/Contents/MacOS/firefox
2216 nate 17 0 3936 115M ? 15.1 1.4 7:14.00 /System/Applications/Messages.app/Contents/MacOS/Messages
1627 nate 17 0 3946 264M ? 10.2 3.2 12:35.00 /Applications/Firefox Nightly.app/Contents/MacOS/plugin-contain
746 nate 17 0 3916 76384 ? 9.9 0.9 1:00.00 /System/Library/PrivateFrameworks/CoreSuggestions.framework/Ver
731 nate 17 0 3926 46880 ? 9.6 0.6 0:14.00 /System/Library/CoreServices/NotificationCenter.app/Contents/Ver
25309 nate 17 0 3916 36240 ? 8.7 0.4 0:02.00 /System/Library/PrivateFrameworks/IMDPersistence.framework/XPC
27255 nate 17 0 3926 22336 ? 7.5 0.3 0:00.00 /usr/sbin/screencapture -pdi -z keyboard.selection
25285 nate 17 0 3916 1872 ? 6.9 0.2 0:05.00 /System/Library/Frameworks/Contacts.framework/Support/contactd
1111 nate 17 0 3946 262M ? 6.5 3.0 0:05.00 /Applications/Firefox Nightly.app/Contents/MacOS/plugin-contain
1097 nate 17 0 3946 324M ? 6.4 4.0 9:26.00 /Applications/Firefox Nightly.app/Contents/MacOS/plugin-contain
25441 nate 17 0 34.66 125M ? 5.7 1.5 0:25.00 /Applications/Spotify.app/Contents/MacOS/Spotify
1129 nate 17 0 3946 228M ? 5.6 2.8 9:24.00 /Applications/Firefox Nightly.app/Contents/MacOS/plugin-contain
1024 nate 17 0 3946 136M ? 4.6 1.7 3:50.00 /Applications/Firefox Nightly.app/Contents/MacOS/plugin-contain
1088 nate 17 0 3916 18656 ? 2.3 0.2 2:56.00 /Applications/Firefox Nightly.app/Contents/MacOS/plugin-contain
25325 nate 17 0 3916 24800 ? 1.5 0.2 0:01.00 /System/Library/PrivateFrameworks/People.framework/People
25386 nate 24 0 3936 40128 ? 1.4 0.5 0:06.00 /System/Applications/Utilities/Terminal.app/Contents/MacOS/Term
1078 nate 17 0 3916 161M ? 1.3 2.0 0:14.00 /System/Library/Services/AppleSpell.service/Contents/MacOS/Apple
789 nate 17 0 3916 29520 ? 1.0 0.4 0:08.00 /System/Library/PrivateFrameworks/IMCore.framework/imagent.app/
25469 nate 8 0 34.20 33772 ? 1.0 0.4 0:03.00 /Applications/Spotify.app/Contents/Frameworks/Spotify Helper.ap
806 nate 17 0 3936 37216 ? 0.9 0.4 0:28.00 /System/Library/Frameworks/CoreServices.framework/Frameworks/Me
27256 nate 17 0 3926 26176 ? 0.9 0.3 0:00.00 /System/Library/CoreServices/screencaptureui.app/Contents/MacOS
718 nate 17 0 3916 31824 ? 0.7 0.4 0:09.00 /System/Library/PrivateFrameworks/IDS.framework/identityservice
666 nate 17 0 3926 44176 ? 0.7 0.5 1:30.00 /System/Library/CoreServices/NotificationCenter.app/Contents/MacOS/C
1103 nate 17 0 3916 25088 ? 0.7 0.3 0:12.00 /System/Library/PrivateFrameworks/PhotoLibraryServices.framework
F1help F2Setup F3Search F4Filter F5Free F6SortBy F7License F8License F9Kill F10Quit
  
```

- ▶ Control output of `top` with your keyboard ([link](#))

File Transfer to VM

- ▶ Option 1: Use `s cp` from command line. Refer to account on VM as `<EID>@ovrw-econ-p02.la.utexas.edu`.
- ▶ Note: This is the address for the IO VM. Other research fields have separate VM allocations; the address will differ slightly.
- ▶ Option 2: Use a GUI `s cp` client like WinSCP.
- ▶ Option 3: RDP local mirror. Access local files in a remote session through your PC's remote desktop GUI. Instructions differ slightly for Windows ([link](#)) and Mac ([link](#)).
- ▶ DO NOT use Box Drive or the Box website. You may be prompted to log into Box Drive upon accessing the VM. Ignore the prompt.

Memory Management

- ▶ Point 1: What is RAM? What is disk memory?
- ▶ Point 2: Size on disk \neq size in RAM
- ▶ Point 3: Multiple processes \implies duplicate memory
- ▶ Conclusion: Be careful! Know your RAM limits!
- ▶ Example: ML packages that bootstrap

Memory Tips I

- ▶ Julia and R: Only pass chunks of your dataframe (if \sim two copies fit in memory)

```
using CSV, DataFrames, Distributed

df = CSV.read("really-big-file.csv") |> DataFrame

pmap(eachrow(df)) do row
    # do things with row...
end
```

- ▶ Python can do similar: use `pandas.iterrows`. Still not copying df to child processes.

Memory Tips II

- ▶ An alternative: *don't* load the file, pass indices to subprocesses and have subprocesses load parts of the file using limit/offset arguments

```
from multiprocessing import Pool
from math import ceil
import pandas as pd
from functools import partial

def process_batch(N, B, batch_num):
    offset = batch_num * B
    limit = B if offset + B < N else N - offset
    df = pd.read_csv("really-big-file.csv", skiprows = offset, nrows = limit)
    # do something with df...

if __name__ == '__main__':
    BATCH_SIZE = 1000
    N_ROWS = # however many rows your CSV has
    f = partial(process_batch, N_ROWS, BATCH_SIZE)
    with Pool(4) as p:
        p.map(f, range(ceil(N_ROWS / BATCH_SIZE)))
```

Memory Tips III

- ▶ Chunked (aka batched) reading is another way to process very large files
 - ▶ `chunksize` in the `pandas.read_csv` function ([link](#))
 - ▶ `read_csv_chunked` in R ([link](#))
 - ▶ `CSV.Chunks` in Julia ([link](#))
 - ▶ `datastore` in Matlab ([link](#))
- ▶ This is about memory management, *not* parallel processing—these functions read one portion of the file at a time. Need to figure out how to add parallelization on top if this is important for performance
- ▶ The Python example code on previous slide is effectively parallel batch processing...

Cluster Access

- ▶ We also have access to one of the world's largest supercomputers at TACC
- ▶ Need a prof to sponsor you as PI and pay for an allocation
- ▶ This is a *cluster* of computers
- ▶ Controlled by Simple Linux Utility for Resource Management (SLURM)
- ▶ For when your computational problems get really big

Pseudo-Random Number Generation I

- ▶ “Random” numbers generated by a computer only approximate a stochastic process. Sequence generated deterministically from an initial seed.
- ▶ How do we seed an RNG so it is reproducible? And how do we do it in parallel?
- ▶ In the initial Python multiprocessing example, we seeded a single RNG that can be passed to functions: best-practice (vs. global RNG) outside of parallel environments.
- ▶ But in parallel, the RNG is passed to each child process so the draws are the same on each child process.
- ▶ Unlikely this is what you want; would be more efficient to draw the random numbers before the parallel process and pass those draws to the child processes.
- ▶ Answer is to spawn a `SeedSequence` and child seeds (equal to number of child processes) from it and RNGs from those, then pass those RNGs to the child processes.
- ▶ More on the hazards and best practices of PRNG in parallel ([link 1](#), [link 2](#)).

Pseudo-Random Number Generation II

- ▶ In Julia, the ability to control how you send variables to subprocesses lets you control whether or not the draws are the same in subprocesses.

```
using Distributed; addprocs(2)

@everywhere seed = randn()
@info "Same" seed_one=@fetchfrom(2, seed) seed_two=@fetchfrom(3, seed)

@info "Different" seeds = pmap(1:5) do i
    i => (procid = myid(), x = randn())
end
```

Github Repository

- ▶ Github repository (link) with all of our code.
- ▶ Includes Python, Matlab, Julia and C implementations of the MNC problem.
- ▶ Also speed comparisons across languages and methods on a binomial choice, Rust (1987)-inspired likelihood calculation.
- ▶ Corresponding Wiki (link), includes some details which expand on these slides.